AASHIMA GUPTA

0413763033 | aashima91.gupta@gmail.com | Melbourne, VIC | LinkedIn | GitHub | Portfolio | Permanent Resident

Operations Data Analyst: Databricks Certified Data Engineer skilled in Power BI, SQL, BigQuery and Fabric, with nearly 4 years building Databricks lakehouse platforms and Azure for banking, payments and enterprise clients.

SELECT ACHIEVEMENTS:

- Data professional with **nearly 4 years in data engineering** and current analytics internship experience, turning raw data into reliable dashboards and insights with SQL, Power BI and BigQuery.
- **Recognition and insights:** Named MVI 1 (Most Valuable Intern) and training MVP at TDAI, showing student success stories converted at **5.6% vs 2.5% CTR** and that only **about 220 of 520** users reaching the payment page paid, leading to a recommendation to improve the payment page.
- **Data quality at scale:** Raised pipeline reliability to **99%+** and cut data quality issues by **40%** at Cognizant, on CDC ingestion handling billions of change events a month.
- **Data modelling:** Designed star schema models with SCD Type 1 and 2 for **15M+ customer profiles and 300M+ transactions** on a payments analytics platform.

WORK EXPERIENCE

Operations Data Analyst (Internship) **Jul 2026 - Present**

THE DATA & AI INSTITUTE (TDAI)
MELBOURNE

Australian data and AI training institute pairing a 12-week analytics program with hands-on internships on real business data.

- Analysed **238 LinkedIn posts** in Excel with formula-driven classification and a weighted scoring model, showing student success stories converted at **5.6% median CTR vs 2.5%** for promotional posts.
- Caught two misleading findings before they reached stakeholders: a post-length effect caused by tagging (17.8% vs 5.0% CTR) and sponsored data overstating the baseline by **10.8%**.
- Built a live Power BI survey dashboard on a Google Sheets star schema model covering **172 responses**, with a 10-theme taxonomy auto-classifying about **97%** of open-text feedback.
- Built a GA4 Power BI dashboard from BigQuery and a Gemini-powered Streamlit web chat, showing organic social converts prospects at **9.7%, 3x paid social and 4x paid search**.

Data Engineer **Feb 2023 - Aug 2023**

COGNIZANT AUSTRALIA
MELBOURNE

Global professional services firm modernising clients in Banking, Insurance and Retail through cloud, data and AI.

- **Configured and optimised Fivetran HVR CDC ingestion into Databricks Bronze, handling millions of change events a month and improving data integrity by 30%.**
- Built Delta Live Tables pipelines with unit tests, validation rules, sanity checks and metadata verification, raising reliability to **99%+** and cutting data quality issues by **40%**.
- Designed Silver models with PK/FK relationships and data masking, and KPI-driven Gold fact and dimension tables for CV compliance, timesheet truancy and utilisation.
- Strengthened production readiness through performance testing, GitHub version control and CI/CD automation, ensuring stable deployments across environments.

Data Engineer **Dec 2021 - Aug 2022**

FULCRUM DIGITAL
PUNE

Digital transformation firm delivering cloud, data platform, automation and AI services to global enterprises.

- Built ETL and ELT pipelines on Azure Data Factory, Databricks and ADLS for a payments analytics platform, running **100+ daily pipelines** through a metadata-driven framework.
- Built a reusable Bronze ingestion framework with schema evolution and reconciliation, cutting new source onboarding from **5 days to 2** and ingestion errors by **25 to 35%**.
- Designed star schema Gold models with SCD Type 1 and 2 for **15M+ customer profiles and 300M+ transactions**, feeding **20+ Power BI reports**.
- Improved analytics query performance by **40 to 60%** through data model and Databricks optimisation.
- Applied data masking and quality rules, securing credentials in Azure Key Vault and deploying through Azure DevOps CI/CD.

Data Engineer **Jul 2019 - Nov 2021**
HCLTECH
PUNE

Global technology company delivering cloud, AI and data transformation across Banking, Insurance and Healthcare.

- Built Azure Databricks lakehouse pipelines for a regulatory reporting platform, ingesting trade data with Auto Loader.
- Processed **150 to 200 GB of trade data a day** through Spark and Delta Lake Medallion pipelines with schema validation, audit logging and lineage.
- Applied **40+ regulatory and business validation rules**, including LEI validation, PII masking and cut duplicate trade records by about **10%**.
- Scheduled and monitored daily Databricks jobs, resolving failures before regulatory reporting deadlines.

System Engineer **Jun 2016 - Jul 2019**
TATA CONSULTANCY SERVICES (TCS)
PUNE

Global IT services leader whose BFSI unit, its largest vertical, serves some of the world's largest financial institutions.

- Developed and maintained Android and Java digital banking applications for BFSI clients within Agile Scrum delivery teams.
- Built secure REST API integrations with token-based authentication between mobile apps and backend banking systems.
- Delivered an insurance proof of concept with a premium calculation engine and risk assessment rules.
- Wrote JUnit and Mockito unit tests and reusable integration components, reducing API integration effort by **30%**.

RELEVANT PORTFOLIO

Insurance Claims Lakehouse [GitHub] I **PERSONAL PROJECT**

- Claims data from 8 JSON and CSV sources (policies, payments, weather, fraud scores) needed integrating to analyse claim speed and risk.
- Built an incremental Medallion pipeline on Databricks with Auto Loader, PySpark, Delta Lake and Unity Catalog, run as a 32-task job deployed through Databricks Asset Bundles.
- Delivered a star schema of 4 dimensions and 4 fact tables with deterministic surrogate keys and hashed PII, guarded by about 55 automated data quality checks.

Fabric Earthquake Analytics [GitHub] **I TRAINING PROJECT**

- USGS earthquake data needed a repeatable daily path from raw GeoJSON to a Power BI report.
- Built Bronze, Silver and Gold PySpark notebooks in a Fabric Lakehouse, run daily by a Data Factory pipeline with dynamic dates.
- Validated 1,420 events in a 6-day test load via the SQL endpoint and published a Power BI map with significance classes.

GA4 Website Analytics: Dashboard and Web Chat I **THE DATA & AI INSTITUTE**

- TDAI had 13 months of GA4 data in BigQuery but no view of which channels drive consultations, enrolments and revenue.
- Modelled sessions and pages into a Power BI star schema within the BigQuery free tier, and built a Streamlit NL-to-SQL web chat powered by Gemini.
- Tracked **$300K+ in website payments** by product and showed over half of visitors leave the payment page without paying.

CORE COMPETENCIES

- **Analytics and BI:** Power BI, DAX, Excel (Power Query, PivotTables), Google Sheets, Google Analytics 4, EDA, Data Visualisation
- **Data Engineering:** Databricks, Microsoft Fabric, Apache Spark, PySpark, Spark SQL, Delta Lake, Delta Live Tables, Auto Loader, Unity Catalog, Fivetran HVR
- **Cloud and Orchestration:** GCP BigQuery, Azure Data Factory, Fabric Data Factory, Azure Data Lake Storage, Azure Key Vault, Databricks Workflows
- **Modelling and Governance:** Star Schema and Dimensional Modelling, Medallion Architecture, SCD Type 1 and 2, Data Quality Frameworks, PII Masking, Data Lineage
- **Languages and DevOps:** SQL, Python, Streamlit, Java, Git, GitHub Actions, Azure DevOps, Databricks Asset Bundles, CI/CD, Agile Scrum

EDUCATION

DATABRICKS CERTIFIED DATA ENGINEER ASSOCIATE I Credential ID 178312810, 2026

THE DATA & AI INSTITUTE I 12-week Data & AI Program with industry internships, MVP award, Jul 2026 - InProgress

MASTER OF COMPUTER APPLICATIONS I Guru Nanak Dev University, Amritsar, 8.1/10 CGPA, 2016